

CSC-591: Foundations of Data Science T/Th. 12:50-2:05pm. EBI-1005.

Ranga **Raju** Vatsavai

Chancellors Faculty Excellence Associate Professor in Geospatial Analytics
Department of Computer Science, North Carolina State University (NCSU)
Associate Director, Center for Geospatial Analytics, NCSU
&
Joint Faculty, Oak Ridge National Laboratory (ORNL)

Administrative

- Any questions

Key points from 8/25

	Attribute Type	Description	Examples	Operations
Categorical Qualitative	Nominal	Nominal attribute values only distinguish. ($=$, \neq)	zip codes, employee ID numbers, eye color, sex: { <i>male</i> , <i>female</i> }	mode, entropy, contingency correlation, χ^2 test
	Ordinal	Ordinal attribute values also order objects. ($<$, $>$)	hardness of minerals, { <i>good</i> , <i>better</i> , <i>best</i> }, grades, street numbers	median, percentiles, rank correlation, run tests, sign tests
Numeric Quantitative	Interval	For interval attributes, differences between values are meaningful. ($+$, $-$)	calendar dates, temperature in Celsius or Fahrenheit	mean, standard deviation, Pearson's correlation, t and F tests
	Ratio	For ratio variables, both differences and ratios are meaningful. ($*$, $/$)	temperature in Kelvin, monetary quantities, counts, age, mass, length, current	geometric mean, harmonic mean, percent variation

Today

- Introduction to Probability

Probability

- Probability is the mathematical language for quantifying uncertainty or randomness
- Given a random **experiment**, a **probability measure** is a **population quantity** that summarizes the randomness

Statistical Experiment

- All statistical experiments have following common things
 - Has more than one outcome
 - Each possible outcome can be specified in advance
 - Outcome depends on chance
- E.g., coin toss, rolling a die

Sample Space and Events

- A **sample space** Ω , is the set of all possible outcomes of an experiment
- Points ω in Ω are called **sample outcomes** or realizations
- **Events** are subsets of Ω

Example (1)

- If we toss a coin twice, then
 $\Omega = \{HH, HT, TH, TT\}$
- Event that the first toss is head
 $A = \{HH, HT\}$

Example (2)

- Let ω be the outcome of measurement of some physical quantity, say temperature, then
 $\Omega = \mathbb{R} = (-\infty, \infty)$
- Event that the measurement is greater than 10, but less than or equal to 20
 $A = (10, 20]$

Basic Set Operations

- Given events A, B, \dots
 - Complement: $A^c = \{\omega \in \Omega; \omega \notin A\}$
 - Union: $A \cup B = \{\omega \in \Omega; \omega \in A \text{ or } \omega \in B \text{ or } \omega \in \text{both}\}$
 - Intersection: $A \cap B = \{\omega \in \Omega; \omega \in A \text{ and } \omega \in B\}$
 - Difference: $A - B = \{\omega : \omega \in A, \omega \notin B\}$
 - Subset: If every element of A is also contained in B , then A is a subset of B $A \subset B$ or, equivalently, $B \supset A$
 - Disjoint: $A_i \cap A_j = \emptyset$
 - A partition of Ω is a sequence of disjoint sets A_1, A_2, \dots such that $\bigcup_{i=1}^{\infty} A_i = \Omega$

Summary of Notations

Table 1. Sample space and events.

Ω	sample space
ω	outcome
A	event (subset of Ω)
$ A $	number of points in A (if A is finite)
A^c	complement of A (not A)
$A \cup B$	union (A or B)
$A \cap B$ or AB	intersection (A and B)
$A - B$	set difference (points in A that are not in B)
$A \subset B$	set inclusion (A is a subset of or equal to B)
\emptyset	null event (always false)
Ω	true event (always true)

Probability

- Probability, $P(A)$, assigns a real number to every event A .
- P is also called a probability distribution or probability measure
- For P to be probability measure, it has to satisfy three axioms

Probability

Definition 2.5 A function \mathbb{P} that assigns a real number $\mathbb{P}(A)$ to each event A is a **probability distribution** or a **probability measure** if it satisfies the following three axioms:

Axiom 1: $\mathbb{P}(A) \geq 0$ for every A

Axiom 2: $\mathbb{P}(\Omega) = 1$

Axiom 3: If A_1, A_2, \dots are disjoint then

$$\mathbb{P}\left(\bigcup_{i=1}^{\infty} A_i\right) = \sum_{i=1}^{\infty} \mathbb{P}(A_i).$$

Many rules can be derived

- If event A implies occurrence of event B , then $\mathbb{P}(A) \leq \mathbb{P}(B)$

Many rules can be derived

- For any two events A and B, the probability that at least one occurs is sum of their probabilities minus their intersection

Group Exercise

- On two tosses of a coin, what is the probability of head on 1st toss or head on 2nd toss (assume all outcomes are equally likely).

Finite Sample Space

- In general, if Ω is finite, and if each outcome is equally likely, then $\mathbb{P}(A) = \frac{|A|}{|\Omega|}$
- $P(A)$ is also called **uniform distribution**
- To compute probabilities, we need to count number of points in A
- Methods for counting points are called combinatorial methods

Counting

- Given n objects, number of ways of ordering these objects is $n! = n(n-1)(n-2) \dots (3)(2)(1)$
- For convenience, we denote $0! = 1$.
- Let us also define $\binom{n}{k} = \frac{n!}{k!(n-k)!}$
- Read as “ n choose k ”, which is the number of distinct ways of choosing k objects from n

Independent Events

- If we flip a fair coin twice, probability of two heads = $\frac{1}{2} \times \frac{1}{2} = \frac{1}{4}$
- Multiplying because we are assuming two events are independent

Definition Two events A and B are **independent** if

$$\mathbb{P}(AB) = \mathbb{P}(A)\mathbb{P}(B)$$

and we write $A \amalg B$. A set of events $\{A_i : i \in I\}$ is independent if

$$\mathbb{P}\left(\bigcap_{i \in J} A_i\right) = \prod_{i \in J} \mathbb{P}(A_i)$$

for every finite subset J of I .

Conditional Probability

Definition If $\mathbb{P}(B) > 0$ then the **conditional probability of A given B** is

$$\mathbb{P}(A|B) = \frac{\mathbb{P}(AB)}{\mathbb{P}(B)}.$$

- Think $\mathbb{P}(A|B)$ as the fraction of times A occurs among those in which B occurs
- Avoid confusion, in general $\mathbb{P}(A|B) \neq \mathbb{P}(B|A)$

Bayes Theorem

Theorem (The Law of Total Probability.) *Let A_1, \dots, A_k be a partition of Ω . Then, for any event B , $\mathbb{P}(B) = \sum_{i=1}^k \mathbb{P}(B|A_i)\mathbb{P}(A_i)$.*

Theorem (Bayes' Theorem.) *Let A_1, \dots, A_k be a partition of Ω such that $\mathbb{P}(A_i) > 0$ for each i . If $\mathbb{P}(B) > 0$ then, for each $i = 1, \dots, k$,*

$$\mathbb{P}(A_i|B) = \frac{\mathbb{P}(B|A_i)\mathbb{P}(A_i)}{\sum_j \mathbb{P}(B|A_j)\mathbb{P}(A_j)}.$$

Example

- Given:
 - A doctor knows that meningitis causes stiff neck 50% of the time
 - Prior probability of any patient having meningitis is $1/50,000$
 - Prior probability of any patient having stiff neck is $1/20$
- If a patient has stiff neck, what's the probability he/she has meningitis?

$$P(M|S) = \frac{P(S|M)P(M)}{P(S)} = \frac{0.5 \times 1/50000}{1/20} = 0.0002$$

Random Variable

- Statistics and DM are concerned with data, so how do we link sample space and events to data?
- A random variable is a **numerical** outcome of an experiment

Definition A random variable is a mapping $X : \Omega \rightarrow \mathbb{R}$ that assigns a real number $X(\omega)$ to each outcome ω .

Random Variable

- A random variable is a **numerical** outcome of an experiment
- From previous lecture, numerical data can be either:
 - Discrete or Continuous
- Discrete: you can count (e.g., #of web hits)
 - We can assign probability to every value they can take
- Continuous: real numbers
 - We can assign probability to ranges

Examples

- Simple examples of discrete r.v.
 - The outcome of the flip of the coin
 - The outcome from the roll of a die
- Complex r.v.
 - Number of vehicles on a road network in a given day (discrete but no upper bound)
 - We can use statistical distribution (e.g., Poisson)

Examples

- Continuous r.v.; i.e., data can take uncountably infinitely values
 - BMI
 - Satellite imagery
 - Temperature
- We use probability distribution (e.g. Gaussian) to assign probabilities

Probability Mass Function (PMF)

- A probability mass function (pmf) is a function that gives the probability that a discrete random variable is exactly equal to some value.

Definition 3.9 X is **discrete** if it takes countably many values

$$\{x_1, x_2, \dots\}.$$

We define the **probability function** or **probability mass function** for X by

$$f_X(x) = \mathbb{P}(X = x).$$

Bernoulli distribution

- Bernoulli distribution is the probability distribution of a random variable which takes value 1 with success probability p and value 0 with failure probability $q=1-p$.
- Coin toss ($H=1$, $T=0$)
- PMF for this distribution can be written as

$$f(k; p) = \begin{cases} p & \text{if } k = 1, \\ 1 - p & \text{if } k = 0. \end{cases}$$

- Can also be written as

$$f(k; p) = p^k (1 - p)^{1-k} \text{ for } k \in \{0, 1\}.$$

R -- Introduction

- Installation
 - <https://www.r-project.org/>
- Manuals/FAQs
- Coursera

Acknowledgements

- Vipin Kumar (Minnesota)
- Jiawei Han (UIUC)
- Hanspeter Pfister (Harvard)

Acknowledgements

- Vipin Kumar (Minnesota)
- Jiawei Han (UIUC)
- Hanspeter Pfister (Harvard)
- Larry Wasserman (CMU)